

Introduction to Research in Computer Science

Lewis University

Spring 2022

Outline

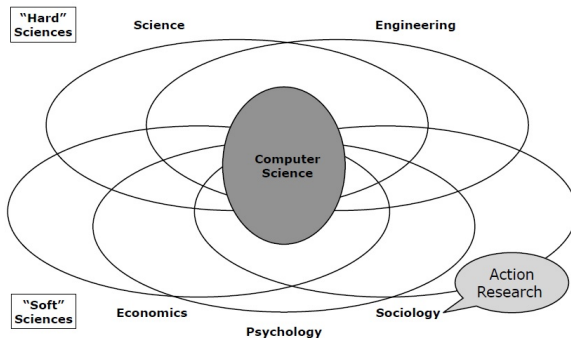
- 1 Course Objectives
- 2 Course Administration
- 3 What is Computer Science?
- 4 Areas of Research in Computer Science
- 5 Algorithms
- 6 Distributed systems
- 7 Building Reliable Software
- 8 Software Verification (Correctness)
- 9 Correctness vs Efficiency

Course Objectives

- Gain an appreciation for the diverse areas encompassing modern computer science
- Understand different approaches to research in computer science, including proofs, observational studies, and implementation-driven research
- Be able to read and ask questions of a computer science research paper
- Present and answer questions about a technical paper
- Explore formal verification as a research area in computer science

- Method of Delivery: Synchronous (live) and Asynchronous (Pre-recording)
- Assessment:
 - Class Contribution (Discussions on BB): 15%
 - Assignments (2 to 3): 35%
 - Course Project (Presentation 15%, Project 35%): 50%
- All assignments and Project (Report and Presentation) should be done in Latex
- Compulsory Readings: Weekly reading, how to read scientific paper links
- Project Proposal's due date: Week #3

What is Computer Science?



What is Computer Science?

- Some combination of:
 - Science
 - Engineering
 - Mathematics
- More precisely
 - “Computer Science” is a name of a field
 - Computer science is a the scientific endeavor relating to computing

Areas of Research in Computer Science

- Abundant-data applications, algorithms, and architectures
- Artificial intelligence and robotics
- Verification, proofs, and automated debugging of hardware designs, software, networking protocols, mathematical theorems, etc
- Bio-informatics and other uses of CS in biology, biomedical engineering, and medicine
- Databases, data centers, information retrieval, and natural-language processing

Areas of Research in Computer Science

- Emerging technologies for computing hardware, communication, and sensing
- Human-computer interaction
- Large-scale networking
- Limits of computation and communication
- Multimedia
- Programming languages and environments
- Security of computer systems and support for digital democracy

A skilled programmer must have good insight into algorithms.

At bachelor level you were offered courses on basic algorithms: searching, sorting, pattern recognition, graph problems, ...

You learned how to detect such subproblems within your programs, and solve them effectively.

You are trained in algorithmic thought for uniprocessor programs (e.g. divide-and-conquer, greedy, memorization).

Consider the following template C code posted on BB.

What algorithm is being implemented here? Write a program in the language of your choice to test this algorithm.

There are three desirable properties for a good algorithm.

What are three desirable properties for a good algorithm?

Designing correct, efficient and implementable algorithms for real-world problem requires access to two distinct bodies of knowledge:

- *Techniques*.....
- *Resources*.....
- **Lesson:** There is a fundamental difference between algorithms, which always produce a correct result, and heuristics, which may usually do a good job but without providing guarantee.

What are three desirable properties for a good algorithm?

Designing correct, efficient and implementable algorithms for real-world problem requires access to two distinct bodies of knowledge:

- *Techniques*.....
- *Resources*.....
- **Lesson:** There is a fundamental difference between algorithms, which always produce a correct result, and heuristics, which may usually do a good job but without providing guarantee.

What are three desirable properties for a good algorithm?

Designing correct, efficient and implementable algorithms for real-world problem requires access to two distinct bodies of knowledge:

- *Techniques*.....
- *Resources*.....
- **Lesson:** There is a fundamental difference between algorithms, which always produce a correct result, and heuristics, which may usually do a good job but without providing guarantee.

Algorithms: Analysis - Big O notation

Complexity measures state how resource consumption (messages, time, space) grows in relation to input size.

For example, if an algorithm has a worst-case message complexity of $O(n^2)$, then for an input of size n , the algorithm in the worst case takes *in the order of* n^2 messages.

Let $f, g : \mathbb{N} \rightarrow \mathbb{R}_{>0}$.

$f = O(g)$ if, for some $C > 0$, $f(n) \leq C \cdot g(n)$ for all $n \in \mathbb{N}$.

$f = \Theta(g)$ if $f = O(g)$ and $g = O(f)$.

Algorithms: Analysis - Big O notation

Complexity measures state how resource consumption (messages, time, space) grows in relation to input size.

For example, if an algorithm has a worst-case message complexity of $O(n^2)$, then for an input of size n , the algorithm in the worst case takes *in the order of* n^2 messages.

Let $f, g : \mathbb{N} \rightarrow \mathbb{R}_{>0}$.

$f = O(g)$ if, for some $C > 0$, $f(n) \leq C \cdot g(n)$ for all $n \in \mathbb{N}$.

$f = \Theta(g)$ if $f = O(g)$ and $g = O(f)$.

Distributed systems

A **distributed system** is an interconnected collection of autonomous processes.

Motivation:

- information exchange
- resource sharing
- parallelization to increase performance
- replication to increase reliability
- multicore programming

Distributed versus uniprocessor

Distributed systems differ from uniprocessor systems in three aspects.

- *Lack of knowledge on the global state*: A process has no up-to-date knowledge on the local states of other processes.

Example: termination and deadlock detection become an issue.

- *Lack of a global time frame*: No total order on events by their temporal occurrence.

Example: mutual exclusion becomes an issue.

- *Nondeterminism*: Execution of processes is nondeterministic, so running a system twice can give different results.

Example: race conditions

Distributed versus uniprocessor

Distributed systems differ from uniprocessor systems in three aspects.

- *Lack of knowledge on the global state*: A process has no up-to-date knowledge on the local states of other processes.

Example: termination and deadlock detection become an issue.

- *Lack of a global time frame*: No total order on events by their temporal occurrence.

Example: mutual exclusion becomes an issue.

- *Nondeterminism*: Execution of processes is nondeterministic, so running a system twice can give different results.

Example: race conditions

Distributed versus uniprocessor

Distributed systems differ from uniprocessor systems in three aspects.


- *Lack of knowledge on the global state*: A process has no up-to-date knowledge on the local states of other processes.
Example: termination and deadlock detection become an issue.
- *Lack of a global time frame*: No total order on events by their temporal occurrence.
Example: mutual exclusion becomes an issue.
- *Nondeterminism*: Execution of processes is nondeterministic, so running a system twice can give different results.
Example: race conditions

Building Reliable Software

- Suppose you run a software company
- Support you've sunk 30+ person-years into developing the “next big thing”:
 - ★ Boeing Dreamliner2 flight controller
 - ★ Autonomous vehicle control software for Tesla
 - ★ Gene therapy DNA tailoring algorithms
 - ★ Super-efficient green-energy power grid controller
- How do you avoid disasters?
 - ★ Turns out software endangers lives

Boeing 737 Max Crashes

- Involved in two crashes
 - ✦ Lion Air Flight 610 on October 29, 2018 — 189 dead
 - ✦ Ethiopian Airlines Flight 302 on March 10, 2019 — 157 dead
- The crash is attributed to design errors including flight control software
 - ✦ The position of larger engines on 737 Max generated additional lift



Engine placement on the third-generation 737 NG (left) versus the MAX (right).

Boeing 737 Max Crashes

- Manoeuvring Characteristics Augmentation System (MCAS)
 - ✦ Software to sense angle of attack (AoA) from a sensor and automatically compensate
- Crashes due to AoA sensor data but also due to MCAS software
- Every time MCAS was switched on and off again, it acted like first time pitching nose lower
 - ✦ incorrect spec not including history
- Max 0.8 degrees pitch during testing, which was changed to 2.4 after
 - ✦ Executing conditions not reflective of testing
- MCAS completely ignored that pilots were desperately pulling back on the yoke
 - ✦ Incorrect spec not considering environment

Not an isolated incident

- NASA's Mars Climate Orbiter
 - ✦ A sub contractor on the engineering team failed to make a simple conversion from English units to metric
 - ✦ \$125 million
- Ariane 5 Flight 501
 - ✦ The software had tried to cram a 64-bit number into a 16-bit space.
 - ✦ Crashed both the primary and the backup computer
 - ✦ \$500 million payload lost + \$XXX to fix the flaw.
- Hawaii Sends Out a State-Wide False Alarm About a Missile Strike
 - ✦ there were “troubling” design flaws in the Hawaii Emergency Management Agency's alert origination software.
- The Equifax social security hack
 - ✦ 143 million of their consumer records (names, SSN, credit card numbers) were stolen by attackers.

Approaches to Validation

- Social
 - ✦ Code reviews
 - ✦ Extreme/pair programming
- Methodological
 - ✦ Design patterns
 - ✦ Test-driven development
 - ✦ Version control
 - ✦ Bug Tracking
- Technological
 - ✦ Static analysis
 - ✦ Fuzzers
- Mathematical
 - ✦ Sound Type Systems
 - ✦ Formal verification



Less formal: Techniques may miss problems in programs

All of these methods should be used!

Even the most formal can still have holes:

- did you prove the right thing?
- do your assumptions match reality?

More formal: eliminate *with certainty* as many problems as possible.

Software: first failure cause of computing systems

Size: from some (tens) of thousands of code lines to some millions of code lines

Development effort:

0,1-0,5 person.year / KLOC (large software)
5-10 person.year / KLOC (critical software)

Share of the effort devoted to fault removal:

45-75%

Fault density:

10-200 faults / KLOC created during development



- static analysis
- proof
- model-checking
- *testing*

0,01-10 faults / KLOC residual in operation



The Quest for Software Correctness

Speech@50-years Celebration CWI Amsterdam

“It is fair to state, that in this digital era correct systems for information processing are more valuable than gold.”



Henk Barendregt

The Importance of Software Correctness

Rapidly increasing **integration of ICT** in different applications

- embedded systems
- communication protocols
- transportation systems

⇒ reliability increasingly depends on software!

Defects can be **fatal** and extremely **costly**

- products subject to mass-production
- safety-critical systems

What is System Verification?

Folklore “definition”

System verification amounts to check whether a system fulfills the qualitative requirements that have been identified

Verification \neq validation

- Verification = “check that we are building the thing **right**”
- Validation = “check that we are building the **right** thing”

Software Verification Techniques

Peer reviewing

- static technique: manual code inspection, no software execution
- detects between 31 and 93% of defects with median of about 60%
- subtle errors (concurrency and algorithm defects) hard to catch


Testing

- dynamic technique in which software is executed

Some figures

- 30% to 50% of software project costs devoted to testing
- more time and effort is spent on validation than on construction
- accepted defect density: about 1 defects per 1,000 code lines

Bug Hunting: the Sooner, the Better



Formal Methods

Intuitive description

Formal methods are the

“applied mathematics for modelling and analysing ICT systems”

Formal methods offer a large potential for:

- obtaining an **early integration** of verification in the design process
- providing **more effective** verification techniques (higher coverage)
- **reducing** the verification time

Usage of formal methods

Highly recommended by IEC, FAA, and NASA for safety-critical software

Formal Verification Techniques for Property P

Deductive methods

- method: provide a formal **proof** that P holds
- tool: theorem prover/proof assistant or proof checker
- applicable if: system has form of a mathematical theory

Model checking

- method: **systematic check** on P in all states
- tool: model checker (SPIN, NUSMV, UPPAAL, ...)
- applicable if: system generates (finite) behavioural model

Model-based simulation or testing

- method: test for P by **exploring possible behaviours**
- applicable if: system defines an executable model

Simulation and Testing

Basic procedure:

- take a model (simulation) or a realisation (testing)
- stimulate it with certain inputs, i.e., the tests
- observe reaction and check whether this is “desired”

Important drawbacks:

- number of possible behaviours is very large (or even infinite)
- unexplored behaviours may contain the fatal bug

About testing ...

testing/simulation can show the presence of errors, **not** their absence

Milestones in Formal Verification

- **Mathematical program correctness** (Turing, 1949)
- **Syntax-based technique for sequential programs** (Hoare, 1969)
 - for a given input, does a computer program generate the correct output?
 - based on compositional proof rules expressed in predicate logic
- **Syntax-based technique for concurrent programs** (Pnueli, 1977)
 - handles properties referring to states during the computation
 - based on proof rules expressed in temporal logic
- **Automated verification of concurrent programs**
 - model-based instead of proof-rule based approach
 - does the concurrent program satisfy a given (logical) property?

Which Approach Should We Consider?

Correctness vs Efficiency: What Comes first?

Are your Undergraduate Knowledge of Algorithms still relevant to the analysis of distributed systems?

Urgent Need to have a **formal framework** for describing Algorithms:
Uniprocessor or distributed systems

In this course, we will analyze **correctness proofs = Formal methods** as the main support for your introduction to research in computer science.

Possible Themes for research initiation: Verification, proofs, and automated debugging of hardware designs, software, networking protocols, formal methods for complex systems, and mathematical theorems.